

# Welcome to Supervised Learning

## Part 1: Introduction to machine learning and the bias-variance tradeoff

Instructor: Andras Zsom

<https://github.com/azsom/Supervised-Learning>  
(<https://github.com/azsom/Supervised-Learning>)

## The topic of the course series: supervised Machine Learning (ML)

- how to build an ML pipeline from beginning to deployment
- we assume you already performed data cleaning
- this is the first course out of 6 courses
  - **Part 1: Introduction to machine learning and the bias-variance tradeoff**
  - Part 2: How to prepare your data for supervised machine learning
  - Part 3: Evaluation metrics in supervised machine learning
  - Part 4: SVMs, Random Forests, XGBoost
  - Part 5: Missing data in supervised ML
  - Part 6: Interpretability
- you can complete the courses in sequence or complete individual courses based on your interest

## Tools

- we use python
  - pros: easy to use for a beginner programmer
  - cons: it is very difficult to write computationally efficient code
  - the divide between users and developers of python packages are wide
- packages we use: sklearn, pandas, numpy, matplotlib, XGBoost, SHAP
- if you are a python user, you need to know exactly what you are doing
  - carefully read the manual, work through the examples, test every line of code you write
    - good test of your understanding: could I write the function/method myself if I had to?
  - do not assume your code works, always test everything
  - there are two types of errors:
    - one that gives an error message
    - usually easy to fix

- the error message tells you in which line the error occurs
- read and understand the error message
- if it's not obvious what the error is, read more on it on stackoverflow for example
- sneaky errors without error message
  - these are tough!
  - your code runs and it gives some output but something is off
  - just staring at the code won't reveal the bug
  - print print print or use a debugger
  - check every line of code, trace issues through the code
- to reduce the number of errors/bugs, do test-driven code development
  - first think about what the output of a function call/cell/piece of a piece of code should be
  - only then write the code
  - check if you got the expected output

## Learning objectives of this course

By the end of the course, you will be able to

- describe how a task like spam filtering can be solved with explicit coding instructions vs. a machine learning algorithm that learns from examples (training data),
- summarize the similarities and differences between supervised and unsupervised ML,
- list the pros and cons of supervised machine learning,
- define the mathematical model behind linear and logistic regression,
- explain what the loss function is,
- describe the two main types of regularization and why it is important,
- perform a simple train/validation/test split on IID data,
- apply linear and logistic regression to datasets,
- tune the regularization hyperparameter,
- identify models with high bias and high variance,
- select the best model and measure its performance on a previously unseen dataset, the test set.

## Module 1: Intro to Machine Learning

### Learning objectives of this module:

- describe how a task like spam filtering can be solved with explicit coding instructions vs. a machine learning algorithm that learns from examples (training data),
- summarize the similarities and differences between supervised and unsupervised ML,
- list the pros and cons of supervised machine learning,

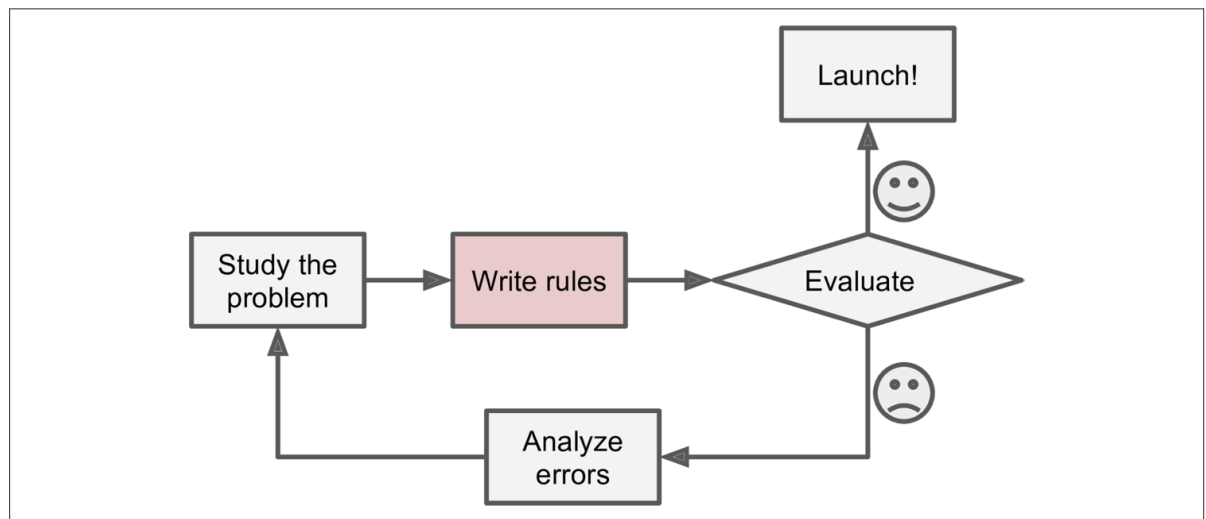
### Supervised ML

- supervised ML is probably the most successful area in ML (based on economic value created)

- **online advertising:** given an ad and user info, will the user click on the ad?
- **real estate:** given home features, can we predict the house price?
- **finance:** given an applicant and a financial product (e.g., a loan), will this applicant be able to successfully pay back the loan?
- **health care:** given a patient, symptoms, and maybe test results, can we predict the illness?
- ...
- supervised ML pros:
  - **automation:** computers perform calculations faster than humans (and computers are cheaper)
  - **learn from examples:** no need to explicitly tell the computer what to do. the computer figures out what to do based on examples (data)
- supervised ML con:
  - it can be difficult or labor-intensive to collect training data
  - there is no guarantee that you will be able to develop an accurate model based on the data you have

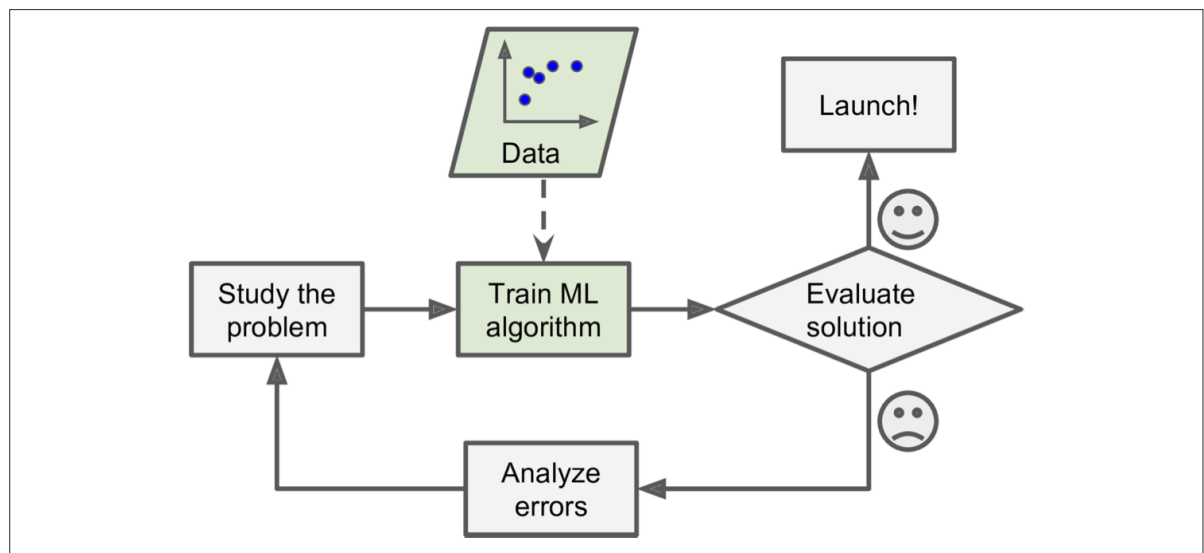
## Example: spam filters

- Traditional coding pipeline with explicit instructions



## Example: spam filters

- ML pipeline



- the data: feature matrix (X) and target variable (Y)
  - X can be structured (tabular data most commonly stored in excel and csv files or SQL databases)
  - X can be unstructured (e.g., images, text, voice recording, video)
  - Y can be categorical, the problem is **classification** (e.g., click or not click on an ad, sick or not sick)
  - Y can be continuous, the problem is **regression** (e.g., predict house price, stock price, age)
  - Y can be missing, the problem is **clustering**
- **we focus on structured data during the course series!**


## Structured data

	X	feature_1	feature_2	...	feature_j	...	feature_m	Y
data_point_1		x_11	x_12	...	x_1j	...	x_1m	y_1
data_point_2		x_21	x_22	...	x_2j	...	x_2m	y_2
...		...	...	...	...	...	...	...
data_point_i		x_i1	x_i2	...	x_ij	...	x_im	y_i
...		...	...	...	...	...	...	...
data_point_n		x_n1	x_n2	...	x_nj	...	x_nm	y_n

## Other areas of ML

- unsupervised ML
  - only the feature matrix X is available, there is no target variable
  - the goal is to find structure (clusters) in the data
  - often used in customer segmentation
- recommender systems
  - recommend products to a customer based on what products similar customers enjoyed

- reinforcement learning
  - the learning system, called an agent, can observe the environment, select and perform actions, and get rewards and penalties in return. Goal: come up with strategy to maximize rewards
  - often used when virtual environment is available (e.g., games like go or warcraft)
  - sounds appealing to use in real environments (like self-driving cars) but agents learn slow, lots of cars would need to be broken to teach an agent to drive this way
- deep learning
  - uses neural networks and often works with unstructured data
  - technically deep learning is supervised or unsupervised
  - extremely successful on large datasets



## Module 2: Overview of linear and logistic regression with regularization

### Learning objectives of this module:

- define the mathematical model behind linear and logistic regression,
- explain what the loss function is,
- describe the two main types of regularization and why it is important,

### Supervised ML algorithms: three parts

- 1) a **mathematical model** ( $f$ ) is used to convert the feature values into a prediction

$f(X_i) = y'_i$ , where  $i$  is the  $i$ th data point in our sample.  $X_i$  is a vector and  $y'_i$  is a number.

- `$f$` is your supervised ML algorithm
- it usually has a number of intrinsic parameters
- 2) an **optimization algorithm** is used to determine the intrinsic parameter values given the training set
  - there are various algorithms
  - e.g., gradient descent, backpropagation
- 3) the optimization algorithm minimizes a metric called **the cost function**
  - the cost function is used to determine the best intrinsic parameters of one model based on the training data

## Linear Regression

```
In [ ]: 1 # these lines are just illustration
        2 # no X_train or y_train are defined yet so it won't run
        3 from sklearn.linear_model import LinearRegression # import the model
        4 LinReg = LinearRegression() # initialize a simple linear regression model
        5 LinReg.fit(X_train,y_train) # we will learn now what happens when you i
```

- This is the **mathematical model**:

$$f(X_i) = y'_i = \theta_0 + X_{i1}\theta_1 + X_{i2}\theta_2 + \dots = \theta_0 + \sum_{j=1}^m \theta_j X_{ij},$$

where  $y'_i$  is the prediction of the linear regression model and  $\theta$  are parameters.

- The **optimization algorithm** is some form of gradient descent
  - we won't go into detail but the basic idea is that gradient descent will find the  $\theta$  values that minimize the cost function on the training data
- The **cost function** is MSE - mean squared error

$$MSE(y, y') = \frac{1}{n} \sum_{i=1}^n (y'_i - y_i)^2$$

## Logistic Regression

```
In [ ]: 1 from sklearn.linear_model import LogisticRegression
        2 LogReg = LogisticRegression() # initialize a simple logistic regression model
        3 LogReg.fit(X_train,y_train) # we will learn now what happens when you issue
```

- name is misleading, logistic regression is for classification problems!
- the model:

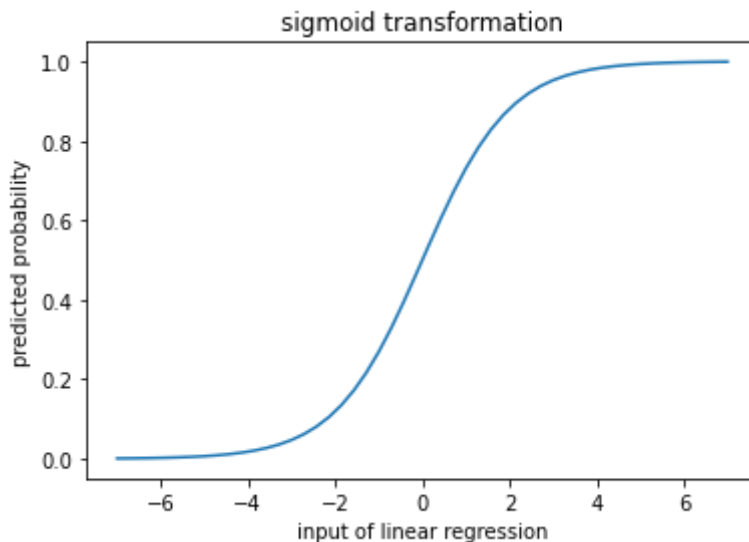
$$f(X_i) = y'_i = \frac{1}{1+e^{-z}}, \text{ where } z = \theta_0 + \sum_{j=1}^m \theta_j x_{ij}$$

- $f(z) = \frac{1}{1+e^{-z}}$  is the sigmoid function which maps real values to be between 0 and 1 such that the real value 0 is mapped to 0.5.

- the output of a sigmoid function can be thought of as a predicted probability.

```
In [15]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 def sigmoid(z):
5     return 1/(1+np.exp(-z))
6
7 z = np.linspace(-7,7,50)
8 print(z)
9
10 plt.plot(z,sigmoid(z))
11 plt.xlabel('input of linear regression')
12 plt.ylabel('predicted probability')
13 plt.title('sigmoid transformation')
14 plt.savefig('figures/sigmoid_trans.png',dpi=300)
15 plt.show()
```

```
[-7.          -6.71428571 -6.42857143 -6.14285714 -5.85714286 -5.57142857
-5.28571429 -5.          -4.71428571 -4.42857143 -4.14285714 -3.85714286
-3.57142857 -3.28571429 -3.          -2.71428571 -2.42857143 -2.14285714
-1.85714286 -1.57142857 -1.28571429 -1.          -0.71428571 -0.42857143
-0.14285714  0.14285714  0.42857143  0.71428571  1.          1.28571429
 1.57142857  1.85714286  2.14285714  2.42857143  2.71428571  3.
 3.28571429  3.57142857  3.85714286  4.14285714  4.42857143  4.71428571
 5.          5.28571429  5.57142857  5.85714286  6.14285714  6.42857143
 6.71428571  7.          ]
```



- The **optimization algorithm** is some form of gradient descent
- the logloss metric is used as a **cost function** in logistic regression

$$L(\theta) = -\frac{1}{N} \sum_{i=1}^n [y_i \ln(y'_i) + (1 - y_i) \ln(1 - y'_i)]$$

- two scenarios:
  - $y_i = 0$  - left term disappears
  - $y_i = 1$  - right term disappears
- $\log(0)$  is undefined
  - $y'_i$  is usually replaced with  $\max(\min(y'_i, 1 - 10^{-15}), 10^{-15})$  to avoid this issue

## The extreme cases

- the classifier is confidently wrong
  - $y'_i = 10^{-15}$  for points in class 1
  - $y'_i = 1 - 10^{-15}$  for points in class 0

$$\logloss = -\frac{1}{N} \sum \ln(10^{-15}) = -\ln(10^{-15})$$
$$\logloss \sim 34.5$$

- the classifier is correct
  - $y'_i = 10^{-15}$  for points in class 0
  - $y'_i = 1 - 10^{-15}$  for points in class 1

$$\logloss = -\frac{1}{N} \sum (1-0)(1 - \ln(1 - 10^{-15})) = 10^{-15} \text{ for class 0}$$

$$\logloss = -\frac{1}{N} \sum 1 * \ln(1 - 10^{-15}) = 10^{-15} \text{ for class 1}$$
$$\logloss \sim 0$$

- the logloss metric also needs to be minimized

## Regularization

- models tend to overfit on the training data and such models don't perform well on previously unseen points
  - a sure sign of overfitting in linear and logistic regression is huge theta values, much larger than the typical ranges of your features and target variable
  - overfitting means that the model fits the noise rather than the underlying structure
    - e.g., fitting a high degree polynomial to a roughly linearly correlated set of points
- one way to address this shortcoming of ML models is regularization
- let's change the cost function and add a **penalty term** for large thetas
- Lasso regression**: regularize using the l1 norm of theta:

$$L(\theta) = \text{original cost} + \frac{\alpha}{m} \sum_{j=0}^m |\theta_j|$$

- Ridge regression**: regularize using the l2 norm of theta:

$$L(\theta) = \text{original cost} + \frac{\alpha}{m} \sum_{j=0}^m \theta_j^2$$

- $\alpha$  is the regularization parameter (0 or larger), it describes how much we penalize large thetas

## Regularization in linear regression

- the original cost function is MSE and we add the penalty term
- Lasso regression**: regularize using the l1 norm of theta:

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n [(\theta_0 + \sum_{j=1}^m \theta_j x_{ij} - y_i)^2] + \frac{\alpha}{m} \sum_{j=0}^m |\theta_j|$$

- Ridge regression**: regularize using the l2 norm of theta:

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n [(\theta_0 + \sum_{j=1}^m \theta_j x_{ij} - y_i)^2] + \frac{\alpha}{m} \sum_{j=0}^m \theta_j^2$$

## Regularization in logistic regression



- the original cost is logloss and we add the penalty term
- **Lasso regression:** regularize using the l1 norm of theta:

$$L(\theta) = -\frac{1}{N} \sum_{i=1}^n [y_i \ln(\frac{1}{1+e^{-\theta_0+\sum_{j=1}^m \theta_j x_{ij}}}) + (1 - y_i) \ln(1 - \frac{1}{1+e^{-\theta_0+\sum_{j=1}^m \theta_j x_{ij}}})] + \frac{\alpha}{m} \sum_{j=0}^m |\theta_j|$$

- **Ridge regression:** regularize using the l2 norm of theta:

$$L(\theta) = -\frac{1}{N} \sum_{i=1}^n [y_i \ln(\frac{1}{1+e^{-\theta_0+\sum_{j=1}^m \theta_j x_{ij}}}) + (1 - y_i) \ln(1 - \frac{1}{1+e^{-\theta_0+\sum_{j=1}^m \theta_j x_{ij}}})] + \frac{\alpha}{m} \sum_{j=0}^m \theta_j^2$$

**Let's translate these concepts to code in the next module!**

## Module 3: The bias-variance tradeoff

### Learning objectives of this module:

- perform a simple train/validation/test split on IID data,
- apply linear and logistic regression to datasets,
- tune the regularization hyperparameter,
- identify models with high bias and high variance,
- select the best model and measure its performance on a previously unseen dataset, the test set.

```
In [16]: 1 # STEP 1: read in the data
2
3 # https://www4.stat.ncsu.edu/~boos/var.select/diabetes.html
4
5 # IID - independent and identically distributed dataset
6
7 import pandas as pd
8 df = pd.read_csv('https://www4.stat.ncsu.edu/~boos/var.select/diabetes.')
9 print(df.head())
10
11 # separate out the feature matrix and the target variable
12 y = df.iloc[:, -1] # the last column is the target variable
13 X = df.iloc[:, :-1] # all but the last column are the features
14 print(y.head())
15 print(X.head())
```

	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6	Y
0	59	2	32.1	101.0	157	93.2	38.0	4.0	4.8598	87	151
1	48	1	21.6	87.0	183	103.2	70.0	3.0	3.8918	69	75
2	72	2	30.5	93.0	156	93.6	41.0	4.0	4.6728	85	141
3	24	1	25.3	84.0	198	131.4	40.0	5.0	4.8903	89	206
4	50	1	23.0	101.0	192	125.4	52.0	4.0	4.2905	80	135

Name: Y, dtype: int64

	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6
0	59	2	32.1	101.0	157	93.2	38.0	4.0	4.8598	87
1	48	1	21.6	87.0	183	103.2	70.0	3.0	3.8918	69
2	72	2	30.5	93.0	156	93.6	41.0	4.0	4.6728	85
3	24	1	25.3	84.0	198	131.4	40.0	5.0	4.8903	89
4	50	1	23.0	101.0	192	125.4	52.0	4.0	4.2905	80

```
In [11]: 1 # STEP 2: split the data
2 from sklearn.model_selection import train_test_split
3
4 X_other, X_test, y_other, y_test = train_test_split(X,y,test_size=0.2,r
5 X_train, X_val, y_train, y_val = train_test_split(X_other,y_other,test_
6
7 # verify the results
8 print(X_train.shape) # 60% for training
9 print(X_val.shape) # 20% for validation
10 print(X_test.shape) # 20% for testing
```

(264, 10)

(89, 10)

(89, 10)

```

In [12]: 1 # STEP 3: preprocess the data
2 from sklearn.preprocessing import StandardScaler
3
4 scaler = StandardScaler() # initialize the scaler
5
6 X_train_prep = scaler.fit_transform(X_train)
7 X_val_prep = scaler.transform(X_val)
8 X_test_prep = scaler.transform(X_test)
9
10 # the _prep objects are now numpy arrays
11 # let's verify that all feature means are 0 and stds are 1
12 print(np.mean(X_train_prep,axis=0))
13 print(np.std(X_train_prep,axis=0))
14 print(np.mean(X_val_prep,axis=0)) # not exactly 0
15 print(np.std(X_val_prep,axis=0)) # not exactly 1
16 print(np.mean(X_test_prep,axis=0)) # not exactly 0
17 print(np.std(X_test_prep,axis=0)) # not exactly 1

[-1.88401483e-16  1.68215610e-16 -1.58122673e-16  2.85966537e-16
 -2.92695161e-16 -3.36431220e-17 -3.70074342e-17  4.47453522e-16
 -7.77156117e-16  4.97918205e-16]
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
[ 0.00482376 -0.03301861  0.12600247 -0.01374034  0.05696159  0.0065957
  0.14555163 -0.13140355 -0.02451996 -0.09520641]
[1.04300891 0.99719443 1.03188936 1.0245732  1.02866879 1.00032617
 0.91058225 0.8268551  1.07662793 1.04397869]
[ 0.0152835  0.05707869 -0.15068117 -0.09269383 -0.07121409 -0.08617701
 0.12650139 -0.15334999 -0.08066075 -0.18983855]
[1.03447167 1.00226924 0.94563471 0.99350518 1.00946002 1.04830518
 1.05149569 0.95477113 0.92799668 0.86719968]


```

```

In [13]: 1 # STEP 4:
          2 # train linear regression models
          3 # tune the regularization parameter
          4 # calculate and visualize train and validation scores
          5 # select the model that performs best on the validation set
          6 # calculate the generalization error using the test set
          7
          8 from sklearn.linear_model import Lasso
          9 from sklearn.metrics import mean_squared_error
         10
         11 alphas = np.logspace(-2,2,13)
         12 print(alphas)
         13
         14 train_scores = []
         15 val_scores = []
         16 models = []
         17 for alpha in alphas:
         18     # initialize the model
         19     linreg = Lasso(alpha=alpha)
         20     # fit it to the training set
         21     linreg.fit(X_train_prep,y_train)
         22     # save the model
         23     models.append(linreg)
         24     # calculate and save train score
         25     y_train_pred = linreg.predict(X_train_prep)
         26     train_score = mean_squared_error(y_train,y_train_pred,squared=False)
         27     train_scores.append(train_score)
         28     # calculate and save val score
         29     y_val_pred = linreg.predict(X_val_prep)
         30     val_score = mean_squared_error(y_val,y_val_pred,squared=False)
         31     val_scores.append(val_score)
         32
         33 # let's visualize the train and validation scores
         34 plt.plot(alphas,train_scores,label='train score')
         35 plt.plot(alphas,val_scores,label='validation score')
         36 plt.xlabel('regularization strength (alpha)',fontsize=13)
         37 plt.ylabel('RMSE',fontsize=13)
         38 plt.semilogx()
         39 plt.legend(fontsize=13)
         40 plt.savefig('figures/bias-variance.png',dpi=300)
         41 plt.show()

[1.00000000e-02 2.15443469e-02 4.64158883e-02 1.00000000e-01
 2.15443469e-01 4.64158883e-01 1.00000000e+00 2.15443469e+00
 4.64158883e+00 1.00000000e+01 2.15443469e+01 4.64158883e+01
 1.00000000e+02]

```



## The bias-variance tradeoff

- high alpha (strong regularization):
  - the model is too simple
  - it performs poorly on both the training and validation sets (RMSEs are large)
  - high bias or low variance model
- low alpha (weak regularization)
  - the model is too complex
  - it performs very well on the training set but it performs comparatively poorly on the validation set
  - low bias or high variance model
- we are looking for the sweet spot in between
  - if your evaluation metric needs to be minimized (e.g., MSE, RMSE, logloss)
    - select the alpha with the smallest validation score
    - the corresponding model is the best
  - if your evaluation metric needs to be maximized (e.g., accuracy, R2)
    - select the alpha with the largest validation score
    - the corresponding model is the best

## Let's select the best model and calculate the generalization error

```
In [14]: 1 indx = np.argmin(val_scores)
2 print('best alpha:', alphas[indx]) # the best alpha value
3 print('best validation score:', val_scores[indx]) # the validation score
4 final_model = models[indx] # pull out the best model
5 y_test_pred = final_model.predict(X_test_prep)
6 gen_error = mean_squared_error(y_test, y_test_pred, squared=False)
7 print('the generalization error:', gen_error) # the error we expect from
```

```
best alpha: 2.154434690031882
best validation score: 57.4873819354221
the generalization error: 54.72060685174691
```

**Things we ignored but will cover in later courses**

## Things we ignored but will cover in later courses

- uncertainty due to splitting
  - different random states in `train_test_split` place different points in train/val/test
  - the random splitting of points introduces an uncertainty in the train/val/test scores and the best alpha values too
  - redo the whole pipeline with a couple of different random states to measure the uncertainty of generalization
- splitting
  - simple IID datasets like this are rare
  - part 2 covers more interesting splitting strategies
- preprocessing
  - this dataset is simple, all features are already numerical so preprocessing was easy (one standard scaler)
  - part 2 deals with more complex preprocessing pipelines
- more complex models have more than one hyperparameters to tune
  - nested for loops should be avoided
  - sklearn has excellent tools to do hyperparameter tuning over an arbitrary number of parameters
  - we cover this in part 4
- interpretability
  - predictions are often not enough!
  - you need to be able to explain how the model works and how it makes predictions
  - the doctor needs to be able to explain to the patient what factors influence the predicted disease progression in general but also for that patient in particular
  - part 6 covers this in detail

In [ ]:

1