

splitting_and_preprocessing

December 23, 2020

#

Welcome to Supervised Learning

##

Part 2: How to prepare your data for supervised machine learning

##

Instructor: Andras Zsom

###

<https://github.com/azsom/Supervised-Learning>

0.1 The topic of the course series: supervised Machine Learning (ML)

- how to build an ML pipeline from beginning to deployment
- we assume you already performed data cleaning
- this is the first course out of 6 courses
 - Part 1: Introduction to machine learning and the bias-variance tradeoff
 - **Part 2: How to prepare your data for supervised machine learning**
 - Part 3: Evaluation metrics in supervised machine learning
 - Part 4: SVMs, Random Forests, XGBoost
 - Part 5: Missing data in supervised ML
 - Part 6: Interpretability
- you can complete the courses in sequence or complete individual courses based on your interest

0.1.1 Structured data

X	feature_1	feature_2	...	feature_j	...	feature_m	Y
data_point_1	x_11	x_12	...	x_1j	...	x_1m	y_1
data_point_2	x_21	x_22	...	x_2j	...	x_2m	y_2
...
data_point_i	x_i1	x_i2	...	x_ij	...	x_im	y_i
...
data_point_n	x_n1	x_n2	...	x_nj	...	x_nm	y_n

We focus on the feature matrix (X) in this course.

0.1.2 Learning objectives of this course

By the end of the course, you will be able to - describe why data splitting is necessary in machine learning - summarize the properties of IID data - list examples of non-IID datasets - apply IID splitting techniques - apply non-IID splitting techniques - identify when a custom splitting strategy is necessary - describe the two motivating concepts behind preprocessing - apply various preprocessors to categorical and continuous features - perform preprocessing with a sklearn pipeline and ColumnTransformer

1 Module 1: Split IID data

1.0.1 Learning objectives of this module:

- describe why data splitting is necessary in machine learning
- summarize the properties of IID data
- apply IID splitting techniques

1.1 Why do we split the data?

- we want to find the best hyper-parameters of our ML algorithms
 - fit models to training data
 - evaluate each model on validation set
 - we find hyper-parameter values that optimize the validation score
- we want to know how the model will perform on previously unseen data - the generalization error
 - apply our final model on the test set

1.1.1 We need to split the data into three parts!

1.2 Ask yourself these questions!

- What is the intended use of the model? What is it supposed to do/predict?
- What data/info do you have available at the time of prediction?
- Your split must mimic the intended use of the model only then will you accurately estimate how well the model will perform on previously unseen points (generalization error).
- two examples:
 - if you want to predict the outcome of a new patient's visit to the ER:
 - * your test score must be based on patients not included in training and validation
 - * your validation score must be based on patients not included in training
 - * points of one patient should not be distributed over multiple sets because your generalization error will be off
 - predict stocks price
 - * it is a time series data
 - * if you predict the stocks price at a certain time in development, make sure that you only use information predating that time

1.3 How should we split the data into train/validation/test?

- data is **Independent and Identically Distributed** (iid)

- all samples stem from the same generative process and the generative process is assumed to have no memory of past generated samples
- identify cats and dogs on images
- predict the house price
- predict if someone's salary is above or below 50k
- examples of not iid data:
 - data generated by time-dependent processes
 - data has group structure (samples collected from e.g., different subjects, experiments, measurement devices)

1.4 Splitting strategies for iid data: basic approach

- 60% train, 20% validation, 20% test for small datasets
- 98% train, 1% validation, 1% test for large datasets
 - if you have 1 million points, you still have 10000 points in validation and test which is plenty to assess model performance

1.4.1 Let's work with the adult data!

<https://archive.ics.uci.edu/ml/datasets/adult>

```
[ ]: import pandas as pd
from sklearn.model_selection import train_test_split

df = pd.read_csv('data/adult_data.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k
    ↳ or less than 50k
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features
print(y)
print(X.head())
```

```
[ ]: help(train_test_split)
```

```
[ ]: random_state = 42

# first split to separate out the training set
X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.
    ↳ 6,random_state=random_state)
print('training set:',X_train.shape, y_train.shape) # 60% of points are in train
print(X_other.shape, y_other.shape) # 40% of points are in other
print(X_train.head())

# second split to separate out the validation and test sets
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.
    ↳ 5,random_state=random_state)
```

```

print('validation set:',X_val.shape, y_val.shape) # 20% of points are in
↳validation
print('test set:',X_test.shape, y_test.shape) # 20% of points are in test
print(X_val.head())
print(X_test.head())

```

1.5 Randomness due to splitting

- the model performance, validation and test scores will change depending on which points are in train, val, test
 - inherent randomness or uncertainty of the ML pipeline
- change the random state a couple of times and repeat the whole ML pipeline to assess how much the random splitting affects your test score
 - you would expect a similar uncertainty when the model is deployed

1.6 Splitting strategies for iid data: k-fold splitting

```

[ ]: from sklearn.model_selection import KFold
help(KFold)

```

```

[ ]: random_state = 42

# first split to separate out the test set
X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.
↳2,random_state=random_state)
print(X_other.shape,y_other.shape)
print('test set:',X_test.shape,y_test.shape)

# do KFold split on other
kf = KFold(n_splits=5,shuffle=True,random_state=random_state)
for train_index, val_index in kf.split(X_other,y_other):
    X_train = X_other.iloc[train_index]
    y_train = y_other.iloc[train_index]
    X_val = X_other.iloc[val_index]
    y_val = y_other.iloc[val_index]
    print('  training set:',X_train.shape, y_train.shape)
    print('  validation set:',X_val.shape, y_val.shape)
    # the validation set contains different points in each iteration
    print(X_val[['age','workclass','education']].head())

```

1.7 How many splits should I create?

- tough question, 3-5 is most common
- if you do n splits, n models will be trained, so the larger the n , the most computationally intensive it will be to train the models
- KFold is usually better suited for small datasets

- KFold is good to estimate uncertainty due to random splitting of train and val, but it is not perfect
 - the test set remains the same

1.7.1 Why shuffling iid data is important?

- by default, data is not shuffled by Kfold which can introduce errors!

1.8 Imbalanced data

- imbalanced data: only a small fraction of the points are in one of the classes, usually ~5% or less but there is no hard limit here
- examples:
 - people visit a bank's website. do they sign up for a new credit card?
 - * most customers just browse and leave the page
 - * usually 1% or less of the customers get a credit card (class 1), the rest leaves the page without signing up (class 0).
 - fraud detection
 - * only a tiny fraction of credit card payments are fraudulent
 - rare disease diagnosis
- the issue with imbalanced data:
 - if you apply train_test_split or KFold, you might not have class 1 points in one of your sets by chance
 - this is what we need to fix

1.9 Solution: stratified splits

```
[ ]: random_state = 42

X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.
↳6,random_state=random_state)
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.
↳5,random_state=random_state)

print('**balance without stratification:**')
# a variation on the order of 1% which would be too much for imbalanced data!
print(y_train.value_counts(normalize=True))
print(y_val.value_counts(normalize=True))
print(y_test.value_counts(normalize=True))

X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.
↳6,stratify=y,random_state=random_state)
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.
↳5,stratify=y_other,random_state=random_state)
print('**balance with stratification:**')
# very little variation (in the 4th decimal point only) which is important if
↳the problem is imbalanced
print(y_train.value_counts(normalize=True))
```

```
print(y_val.value_counts(normalize=True))
print(y_test.value_counts(normalize=True))
```

1.10 Stratified folds

```
[ ]: from sklearn.model_selection import StratifiedKFold
help(StratifiedKFold)
```

```
[ ]: # what we did before: variance in balance on the order of 1%
random_state = 42

X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.
↳2,random_state=random_state)
print('test balance:',y_test.value_counts(normalize=True))

# do KFold split on other
kf = KFold(n_splits=5,shuffle=True,random_state=random_state)
for train_index, val_index in kf.split(X_other,y_other):
    X_train = X_other.iloc[train_index]
    y_train = y_other.iloc[train_index]
    X_val = X_other.iloc[val_index]
    y_val = y_other.iloc[val_index]
    print('train balance:')
    print(y_train.value_counts(normalize=True))
    print('val balance:')
    print(y_val.value_counts(normalize=True))
```

```
[ ]: # stratified K Fold: variation in balance is very small (4th decimal point)
random_state = 42

# stratified train-test split
X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.
↳2,stratify=y,random_state=random_state)
print('test balance:',y_test.value_counts(normalize=True))

# do StratifiedKFold split on other
kf = StratifiedKFold(n_splits=5,shuffle=True,random_state=random_state)
for train_index, val_index in kf.split(X_other,y_other):
    X_train = X_other.iloc[train_index]
    y_train = y_other.iloc[train_index]
    X_val = X_other.iloc[val_index]
    y_val = y_other.iloc[val_index]
    print('train balance:')
    print(y_train.value_counts(normalize=True))
    print('val balance:')
    print(y_val.value_counts(normalize=True))
```

2 Module 2: Split non-IID data

2.0.1 Learning objectives of this module:

- list examples of non-IID datasets
- apply non-IID splitting techniques
- identify when a custom splitting strategy is necessary

2.1 Examples of non-iid data

- if there is any sort of time or group structure in your data, it is likely non-iid
 - group structure:
 - * each point is someone's visit to the ER and some people visited the ER multiple times
 - * each point is a customer's visit to website and customers tend to return regularly
 - time structure
 - * each point is the stocks price at a given time
 - * each point is a person's health or activity status

2.2 Group-based split: GroupShuffleSplit

```
[ ]: import numpy as np
from sklearn.model_selection import GroupShuffleSplit
X = np.ones(shape=(8, 2))
y = np.ones(shape=(8, 1))
groups = np.array([1, 1, 2, 2, 2, 3, 3, 3])

gss = GroupShuffleSplit(n_splits=10, train_size=.8, random_state=42)

for train_idx, test_idx in gss.split(X, y, groups):
    print("TRAIN:", train_idx, "TEST:", test_idx)
```

2.3 Group-based split: GroupKFold

```
[ ]: from sklearn.model_selection import GroupKFold

group_kfold = GroupKFold(n_splits=3)

for train_index, test_index in group_kfold.split(X, y, groups):
    print("TRAIN:", train_index, "TEST:", test_index)
```

```
[ ]: help(GroupKFold)
```

2.4 Data leakage in time series data is similar!

- do NOT use information in validation or test which will not be available once your model is deployed
 - don't use future information!

```
[ ]: import numpy as np
from sklearn.model_selection import TimeSeriesSplit
X = np.array([[1, 2], [3, 4], [1, 2], [3, 4], [1, 2], [3, 4]])
y = np.array([1, 2, 3, 4, 5, 6])
tscv = TimeSeriesSplit()
for train_index, test_index in tscv.split(X):
    print("TRAIN:", train_index, "TEST:", test_index)
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
```

2.5 When should you develop your own splitting function?

- there are certain splitting strategies sklearn can't handle at the moment
 - time series data with group structure is one example
 - if you want certain groups to be in certain sets
 - * group structure in classification where all points in a group belong to a certain class
 - * you might want a roughly equal number of groups of each class to be in each set
- check out the [model selection](#) part of sklearn
 - if the splitting strategy you want to follow is not there, implement your own function

3 Module 3: Preprocess continuous and categorical features

3.0.1 Learning objectives of this module:

- describe the two motivating concepts behind preprocessing
- apply various preprocessors to categorical and continuous features
- perform preprocessing with a sklearn pipeline and ColumnTransformer

3.0.2 Data almost never comes in a format that's directly usable in ML

- ML works with numerical data but some columns are text (e.g., home country, educational level, gender, race)
 - some ML algorithms accept (and prefer) a non-numerical feature matrix (like [CatBoost](#)) but that's not standard
 - sklearn throws an error message if the feature matrix contains non-numerical elements
- the order of magnitude of numerical features can vary greatly which is not good for most ML algorithms (e.g., salary in USD, age in years, time spent on the site in sec)
 - many ML algorithms are distance-based and they perform better and converge faster if the features are standardized (features have a mean of 0 and the same standard deviation, usually 1)
 - * Lasso and Ridge regression because of the penalty term, K Nearest Neighbors, SVM, linear models if you want to use the coefficients to measure feature importance (more on this in part 6), neural networks
 - tree-based methods don't require standardization
 - check out part 1 to learn more about linear and logistic regression, Lasso and Ridge
 - check out part 4 to learn more about SVMs, tree-based methods, and K Nearest Neighbors

3.0.3 scikit-learn transformers to the rescue!

Preprocessing is done with various transformers. All transformers have three methods: - **fit** method: estimates parameters necessary to do the transformation, - **transform** method: transforms the data based on the estimated parameters, - **fit_transform** method: both steps are performed at once, this can be faster than doing the steps separately.

3.0.4 Transformers we cover

- **OrdinalEncoder** - converts categorical features into an integer array
- **OneHotEncoder** - converts categorical features into dummy arrays
- **StandardScaler** - standardizes continuous features by removing the mean and scaling to unit variance

3.1 Ordered categorical data: OrdinalEncoder

Let's assume we have a categorical feature and training and test sets

The categories can be ordered or ranked

E.g., educational level in the adult dataset

```
[ ]: import pandas as pd

train_edu = {'educational level':
    ↳ ['Bachelors', 'Masters', 'Bachelors', 'Doctorate', 'HS-grad', 'Masters']}
test_edu = {'educational level':
    ↳ ['HS-grad', 'Masters', 'Masters', 'College', 'Bachelors']}

X_train = pd.DataFrame(train_edu)
X_test = pd.DataFrame(test_edu)

[ ]: from sklearn.preprocessing import OrdinalEncoder
help(OrdinalEncoder)

[ ]: # initialize the encoder
cats = ['HS-grad', 'Bachelors', 'Masters', 'Doctorate']

enc = OrdinalEncoder(categories = [cats]) # The ordered list of
# categories need to be provided. By default, the categories are alphabetically
↳ ordered!

# fit the training data
enc.fit(X_train)
# print the categories - not really important because we manually gave the
↳ ordered list of categories
print(enc.categories_)
# transform X_train. We could have used enc.fit_transform(X_train) to combine
↳ fit and transform
```

```

X_train_oe = enc.transform(X_train)
print(X_train_oe)
# transform X_test
X_test_oe = enc.transform(X_test) # OrdinalEncoder always throws an error
    ↳ message if
                                   # it encounters an unknown category in test
print(X_test_oe)

```

3.2 Unordered categorical data: one-hot encoder

some categories cannot be ordered. e.g., workclass, relationship status

first feature: gender (male, female, unknown)

second feature: browser used

these categories cannot be ordered

```

[ ]: train = {'gender': ['Male', 'Female', 'Unknown', 'Male', 'Female', 'Female'], \
            'browser': ['Safari', 'Safari', 'Internet_
    ↳ Explorer', 'Chrome', 'Chrome', 'Internet Explorer']}
test = {'gender': ['Female', 'Male', 'Unknown', 'Female'], 'browser':
    ↳ ['Chrome', 'Firefox', 'Internet Explorer', 'Safari']}

X_train = pd.DataFrame(train)
X_test = pd.DataFrame(test)

```

```

[ ]: # How do we convert this to numerical features?
from sklearn.preprocessing import OneHotEncoder

help(OneHotEncoder)

```

```

[ ]: # initialize the encoder
enc = OneHotEncoder(sparse=False) # by default, OneHotEncoder returns a sparse
    ↳ matrix. sparse=False returns a 2D array
# fit the training data
enc.fit(X_train)
print('categories:', enc.categories_)
print('feature names:', enc.get_feature_names())
# transform X_train
X_train_ohe = enc.transform(X_train)
#print(X_train_ohe)
# do all of this in one step
X_train_ohe = enc.fit_transform(X_train)
print(X_train_ohe)

# transform X_test
X_test_ohe = enc.transform(X_test)

```

```
print('X_test transformed')
print(X_test_ohc)
```

3.3 Continuous features: StandardScaler

```
[ ]: train = {'salary': [50_000, 75_000, 40_000, 1_000_000, 30_000, 250_000, 35_000, 45_000]}
      test = {'salary': [25_000, 55_000, 1_500_000, 60_000]}
```

```
X_train = pd.DataFrame(train)
X_test = pd.DataFrame(test)
```

```
[ ]: from sklearn.preprocessing import StandardScaler
      help(StandardScaler)
```

```
[ ]: scaler = StandardScaler()
      print(scaler.fit_transform(X_train))
      print(scaler.transform(X_test))
```

3.4 How and when to do preprocessing in the ML pipeline?

- **SPLIT YOUR DATA FIRST!**
- **APPLY TRANSFORMER.FIT ONLY ON YOUR TRAINING DATA!** Then transform the validation and test sets.
- One of the most common mistake practitioners make is leaking statistics!
 - fit_transform is applied to the whole dataset, then the data is split into train/validation/test
 - * this is wrong because the test set statistics impacts how the training and validation sets are transformed
 - * but the test set must be separated from train and val, and val must be separated from train
 - or fit_transform is applied to the train, then fit_transform is applied to the validation set, and fit_transform is applied to the test set
 - * this is wrong because the relative position of the points change

3.5 Scikit-learn's pipelines

- Preprocessing and model training (not the splitting) can be chained together into a scikit-learn pipeline which consists of transformers and one final estimator which is usually your classifier or regression model.
- It neatly combines the preprocessing steps and it helps to avoid leaking statistics.

https://scikit-learn.org/stable/auto_examples/compose/plot_column_transformer_mixed_types.html

```
[ ]: import pandas as pd
      import numpy as np

      from sklearn.compose import ColumnTransformer
      from sklearn.pipeline import Pipeline
```

```

from sklearn.preprocessing import StandardScaler, OneHotEncoder,
↳OrdinalEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split

np.random.seed(0)

df = pd.read_csv('data/adult_data.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k
↳or less than 50k
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features

random_state = 42

# first split to separate out the training set
X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.
↳6,random_state=random_state)

# second split to separate out the validation and test sets
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.
↳5,random_state=random_state)

```

```

[ ]: # collect which encoder to use on each feature
# needs to be done manually
ordinal_ftrs = ['education']
ordinal_cats = [[' Preschool', ' 1st-4th', ' 5th-6th', ' 7th-8th', ' 9th', ' 10th', '
↳11th', ' 12th', ' HS-grad', \
                ' Some-college', ' Assoc-voc', ' Assoc-acdm', ' Bachelors', '
↳Masters', ' Prof-school', ' Doctorate']]
onehot_ftrs =
↳['workclass', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'native-country']
std_ftrs = ['capital-gain', 'capital-loss', 'age', 'hours-per-week']

# collect all the encoders
preprocessor = ColumnTransformer(
    transformers=[
        ('ord', OrdinalEncoder(categories = ordinal_cats), ordinal_ftrs),
        ('onehot', OneHotEncoder(sparse=False,handle_unknown='ignore'),
↳onehot_ftrs),
        ('std', StandardScaler(), std_ftrs)])

# for now we only preprocess, later on we will add other steps here
# note the final scaler which is a standard scaler
# the ordinal and one hot encoded features do not have a mean of 0 and an std
↳of 1

```

```

# the final scaler standardizes those features
clf = Pipeline(steps=[('preprocessor', preprocessor), ('final_
→scaler', StandardScaler())])

X_train_prep = clf.fit_transform(X_train)
X_val_prep = clf.transform(X_val)
X_test_prep = clf.transform(X_test)

print(X_train.shape)
print(X_train_prep.shape)

print(np.mean(X_train_prep,axis=0))
print(np.std(X_train_prep,axis=0))
print(np.mean(X_val_prep,axis=0))
print(np.std(X_val_prep,axis=0))
print(np.mean(X_test_prep,axis=0))
print(np.std(X_test_prep,axis=0))

```

[]: