

# interpretability

January 18, 2021

#

Welcome to Supervised Learning

##

Part 6: Interpretability

##

Instructor: Andras Zsom

###

<https://github.com/azsom/Supervised-Learning>

## 0.1 The topic of the course series: supervised Machine Learning (ML)

- how to build an ML pipeline from beginning to deployment
- we assume you already performed data cleaning
- this is the sixth course out of 6 courses
  - Part 1: Introduction to machine learning and the bias-variance tradeoff
  - Part 2: How to prepare your data for supervised machine learning
  - Part 3: Evaluation metrics in supervised machine learning
  - Part 4: Non-linear supervised machine learning algorithms
  - Part 5: Missing data in supervised ML
  - **Part 6: Interpretability**
- you can complete the courses in sequence or complete individual courses based on your interest

### 0.1.1 Structured data

X	feature_1	feature_2	...	feature_j	...	feature_m	Y
<b>data_point_1</b>	x_11	x_12	...	x_1j	...	x_1m	<b>y_1</b>
<b>data_point_2</b>	x_21	x_22	...	x_2j	...	x_2m	<b>y_2</b>
...	...	...	...	...	...	...	...
<b>data_point_i</b>	x_i1	x_i2	...	x_ij	...	x_im	<b>y_i</b>
...	...	...	...	...	...	...	...
<b>data_point_n</b>	x_n1	x_n2	...	x_nj	...	x_nm	<b>y_n</b>

### 0.1.2 Learning objectives of this course

By the end of the course, you will be able to - Summarize why it is important to explain models - Describe why additional tools are necessary to explain non-linear models - Review the difference between global and local feature importance metrics - Use the coefficients of linear models to measure feature importance - Apply permutation feature importance to calculate global feature importances - Describe some model-specific approaches to measure global feature importance - Describe the intuition behind SHAP values - Create force, dependence, and summary plots to aid local interpretability

## 1 Module 1: Global feature importance metrics in linear models

### 1.0.1 Learning objectives of this module:

- Summarize why it is important to explain models
- Describe why additional tools are necessary to explain non-linear models
- Review the difference between global and local feature importance metrics
- Use the coefficients of linear models to measure feature importance

### 1.1 Motivation

- debugging ML models is tough
  - a model that runs without errors/warning is not necessarily correct
- how do you know that your model is correct?
  - check test set predictions
    - \* in regression: check points with a large difference between true and predicted values
    - \* in classification: confusion matrix, check out FPs and FNs
  - inspect your model
    - \* especially useful for non-linear models
    - \* metrics to measure how much a model depends on a feature is one way to inspect your model

#### 1.1.1 Motivation

- local feature importance improves the interpretability of complex models
- check out [this page](#) for a good example

#### 1.1.2 Motivation

- can we trust the model?
  - global feature importance: does the model make predictions based on reasonable features?
  - local feature importance: can we trust the model's prediction for one specific data point?
- global feature importance is often not enough especially when you work with human data
  - medical: the doctor needs to be able to explain the reasoning behind the model prediction to the patient
  - finance: customer wants to know why they were declined a loan/mortgage/credit card/etc

## 1.2 Coefficients of linear models

- the coefficients of linear and logistic regression can be used as a measure of feature importance  
**ONLY IF** all features have a zero mean and the same standard deviation (usually 1)
  - all features meaning that the one-hot encoded and ordinal features as well!
- then the absolute value of the coefficients can be used to rank them

### 1.2.1 Let's work with the adult dataset

- <https://archive.ics.uci.edu/ml/datasets/adult>

```
[1]: import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.svm import SVC
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
import matplotlib.pyplot as plt

df = pd.read_csv('data/adult_data.csv')
label = 'gross-income'
y = LabelEncoder().fit_transform(df[label])
df.drop(columns=[label], inplace=True)
X = df
ftr_names = X.columns
print(X.head())
print(y)
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country
--	--------------	--------------	----------------	----------------

0	2174	0	40	United-States
1	0	0	13	United-States
2	0	0	40	United-States
3	0	0	40	United-States
4	0	0	40	Cuba

[0 0 0 ... 0 0 1]

```
[2]: from sklearn.linear_model import LogisticRegression
def ML_pipeline_kfold_LR1(X,y,random_state,n_folds):
    # create a test set
    X_other, X_test, y_other, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state = random_state)
    # splitter for _other
    kf =
    ↪StratifiedKFold(n_splits=n_folds,shuffle=True,random_state=random_state)
    # create the pipeline: preprocessor + supervised ML method
    cat_ftrs =
    ↪['workclass','education','marital-status','occupation','relationship','race','sex','native-
    cont_ftrs =
    ↪['age','fnlwgt','education-num','capital-gain','capital-loss','hours-per-week']
    # one-hot encoder
    categorical_transformer = Pipeline(steps=[
        ('onehot', OneHotEncoder(sparse=False,handle_unknown='ignore'))])
    # standard scaler
    numeric_transformer = Pipeline(steps=[
        ('scaler', StandardScaler())])
    preprocessor = ColumnTransformer(
        transformers=[
            ('num', numeric_transformer, cont_ftrs),
            ('cat', categorical_transformer, cat_ftrs)])
    pipe =
    ↪make_pipeline(preprocessor,LogisticRegression(penalty='l2',solver='lbfgs'))
    # the parameter(s) we want to tune
    param_grid = {'logisticregression__C': [0.01, 0.1, 1, 10,100]}
    # prepare gridsearch
    grid = GridSearchCV(pipe, param_grid=param_grid,cv=kf, return_train_score =
    ↪True,n_jobs=-1)
    # do kfold CV on _other
    grid.fit(X_other, y_other)
    feature_names = cont_ftrs + \
        list(grid.best_estimator_[0].named_transformers_['cat'][0].
    ↪get_feature_names(cat_ftrs))
    return grid, np.array(feature_names), X_test, y_test

[3]: grid, feature_names, X_test, y_test = ML_pipeline_kfold_LR1(X,y,42,4)
print('test score:',grid.score(X_test,y_test))
coefs = grid.best_estimator_[-1].coef_[0]
```

```
sorted_indcs = np.argsort(np.abs(coefs))

plt.rcParams.update({'font.size': 14})
plt.barh(np.arange(10), coefs[sorted_indcs[-10:]]))
plt.yticks(np.arange(10), feature_names[sorted_indcs[-10:]])
plt.xlabel('coefficient')
plt.title('not all scaled')
plt.tight_layout()
plt.savefig('figures/LR_coefs_notscaled.png', dpi=300)
plt.show()
```

/Users/azsom/opt/anaconda3/envs/data1030/lib/python3.7/site-packages/sklearn/linear\_model/\_logistic.py:764: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:


<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

test score: 0.8581298940580377




```
[4]: from sklearn.linear_model import LogisticRegression
def ML_pipeline_kfold_LR2(X,y,random_state,n_folds):
    # create a test set
    X_other, X_test, y_other, y_test = train_test_split(X, y, test_size=0.2,
    ↳random_state = random_state)
    # splitter for _other
    kf =
    ↳StratifiedKfold(n_splits=n_folds,shuffle=True,random_state=random_state)
    # create the pipeline: preprocessor + supervised ML method
    cat_ftrs =
    ↳['workclass','education','marital-status','occupation','relationship','race','sex','native-
    cont_ftrs =
    ↳['age','fnlwgt','education-num','capital-gain','capital-loss','hours-per-week']
    # one-hot encoder
    categorical_transformer = Pipeline(steps=[
        ('onehot', OneHotEncoder(sparse=False,handle_unknown='ignore'))])
    # standard scaler
    numeric_transformer = Pipeline(steps=[
        ('scaler', StandardScaler())])
    preprocessor = ColumnTransformer(
        transformers=[
            ('num', numeric_transformer, cont_ftrs),
            ('cat', categorical_transformer, cat_ftrs)])
    final_scaler = StandardScaler()
    pipe =
    ↳make_pipeline(preprocessor,final_scaler,LogisticRegression(penalty='l2',solver='lbfgs'))
    # the parameter(s) we want to tune
    param_grid = {'logisticregression__C': [0.01, 0.1, 1, 10,100]}
    # prepare gridsearch
    grid = GridSearchCV(pipe, param_grid=param_grid,cv=kf, return_train_score =
    ↳True,n_jobs=-1)
    # do kfold CV on _other
    grid.fit(X_other, y_other)
    feature_names = cont_ftrs + \
        list(grid.best_estimator_[0].named_transformers_['cat'][0].
    ↳get_feature_names(cat_ftrs))
    return grid, np.array(feature_names), X_test, y_test
```

```
[5]: grid, feature_names, X_test, y_test = ML_pipeline_kfold_LR2(X,y,42,4)
print('test score:',grid.score(X_test,y_test))
coefs = grid.best_estimator_[-1].coef_[0]
sorted_indcs = np.argsort(np.abs(coefs))

plt.rcParams.update({'font.size': 14})
plt.barh(np.arange(10),coefs[sorted_indcs[-10:]])
plt.yticks(np.arange(10),feature_names[sorted_indcs[-10:]])
```

```
plt.xlabel('coefficient')
plt.title('all scaled')
plt.tight_layout()
plt.savefig('figures/LR_coefs_scaled.png',dpi=300)
plt.show()
```

test score: 0.857976354982343



## 2 Module 2: Global feature importance metrics in non-linear models

### 2.0.1 Learning objectives of this module:

- Apply permutation feature importance to calculate global feature importances
- Describe some model-specific approaches to measure global feature importance

### 2.1 Permutation feature importance

- model agnostic, you can use it with any supervised ML model
- steps:
  - train a model and calculate a test score :)
  - randomly shuffle a single feature in the test set
  - recalculate the test score with the shuffled data

- model score worsens because the shuffling breaks the relationship between feature and target
- the larger the difference, the more important the feature is

```
[6]: import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.svm import SVC
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
import matplotlib.pyplot as plt

df = pd.read_csv('data/adult_data.csv')
label = 'gross-income'
y = LabelEncoder().fit_transform(df[label])
df.drop(columns=[label], inplace=True)
X = df
ftr_names = X.columns
print(X.head())
print(y)
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country
0	2174	0	40	United-States
1	0	0	13	United-States
2	0	0	40	United-States
3	0	0	40	United-States
4	0	0	40	Cuba

[0 0 0 ... 0 0 1]



```
[7]: def ML_pipeline_kfold(X,y,random_state,n_folds):
    # create a test set
    X_other, X_test, y_other, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state = random_state)
    # splitter for _other
    kf =
    ↪StratifiedKfold(n_splits=n_folds,shuffle=True,random_state=random_state)
    # create the pipeline: preprocessor + supervised ML method
    cat_ftrs =
    ↪['workclass','education','marital-status','occupation','relationship','race','sex','native-
    cont_ftrs =
    ↪['age','fnlwgt','education-num','capital-gain','capital-loss','hours-per-week']
    # one-hot encoder
    categorical_transformer = Pipeline(steps=[
        ('onehot', OneHotEncoder(sparse=False,handle_unknown='ignore'))])
    # standard scaler
    numeric_transformer = Pipeline(steps=[
        ('scaler', StandardScaler())])
    preprocessor = ColumnTransformer(
        transformers=[
            ('num', numeric_transformer, cont_ftrs),
            ('cat', categorical_transformer, cat_ftrs)])
    pipe = make_pipeline(preprocessor,SVC())
    # the parameter(s) we want to tune
    param_grid = {'svc__C': [0.01, 0.1, 1, 10, 100],
        'svc__gamma': [0.01, 0.1, 1, 10, 100]}
    # prepare gridsearch
    grid = GridSearchCV(pipe, param_grid=param_grid,cv=kf, return_train_score =
    ↪True,n_jobs=-1,verbose=10)
    # do kfold CV on _other
    grid.fit(X_other, y_other)
    return grid, X_test, y_test
```

### 2.1.1 Be careful, SVM is used on a relatively large dataset

```
[8]: grid, X_test, y_test = ML_pipeline_kfold(X,y,42,4)
print(grid.best_score_)
print(grid.score(X_test,y_test))
print(grid.best_params_)

# save the output so I can use it later
import pickle
file = open('results/grid.save', 'wb')
pickle.dump((grid,X_test,y_test),file)
file.close()
```

Fitting 4 folds for each of 25 candidates, totalling 100 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done   2 tasks      | elapsed:   1.5min
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:   4.9min
[Parallel(n_jobs=-1)]: Done  16 tasks      | elapsed:   6.0min
[Parallel(n_jobs=-1)]: Done  25 tasks      | elapsed:   9.6min
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:  15.4min
[Parallel(n_jobs=-1)]: Done  45 tasks      | elapsed:  18.6min
[Parallel(n_jobs=-1)]: Done  56 tasks      | elapsed:  25.2min
[Parallel(n_jobs=-1)]: Done  69 tasks      | elapsed:  33.2min
[Parallel(n_jobs=-1)]: Done  82 tasks      | elapsed:  41.1min
[Parallel(n_jobs=-1)]: Done  96 out of 100 | elapsed: 51.4min remaining:  2.1min
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 55.6min finished
```

0.8545377764127764

0.8624289881774911

{'svc\_\_C': 1, 'svc\_\_gamma': 0.1}

```
[9]: import pickle
file = open('results/grid.save', 'rb')
grid, X_test, y_test = pickle.load(file)
file.close()

np.random.seed(42)

nr_runs = 10
scores = np.zeros([len(ftr_names),nr_runs])

test_score = grid.score(X_test,y_test)
print('test score = ',test_score)
print('test baseline = ',np.sum(y_test == 0)/len(y_test))
# loop through the features
for i in range(len(ftr_names)):
    print('shuffling '+str(ftr_names[i]))
    acc_scores = []
    for j in range(nr_runs):
        X_test_shuffled = X_test.copy()
        X_test_shuffled[ftr_names[i]] = np.random.
        ↳ permutation(X_test[ftr_names[i]].values)
        acc_scores.append(grid.score(X_test_shuffled,y_test))
    print('    shuffled test score:',np.around(np.mean(acc_scores),3),'+/-',np.
    ↳ around(np.std(acc_scores),3))
    scores[i] = acc_scores
```

test score = 0.8624289881774911

test baseline = 0.7587901120835252

shuffling age

shuffled test score: 0.851 +/- 0.002

shuffling workclass

```


    shuffled test score: 0.861 +/- 0.001
shuffling fnlwgt
    shuffled test score: 0.862 +/- 0.001
shuffling education
    shuffled test score: 0.86 +/- 0.001
shuffling education-num
    shuffled test score: 0.839 +/- 0.002
shuffling marital-status
    shuffled test score: 0.842 +/- 0.002
shuffling occupation
    shuffled test score: 0.844 +/- 0.002
shuffling relationship
    shuffled test score: 0.851 +/- 0.003
shuffling race
    shuffled test score: 0.862 +/- 0.0
shuffling sex
    shuffled test score: 0.862 +/- 0.0
shuffling capital-gain
    shuffled test score: 0.823 +/- 0.001
shuffling capital-loss
    shuffled test score: 0.855 +/- 0.001
shuffling hours-per-week
    shuffled test score: 0.855 +/- 0.002
shuffling native-country
    shuffled test score: 0.862 +/- 0.001

```

```

[10]: sorted_indcs = np.argsort(np.mean(scores,axis=1))[:-1]
plt.rcParams.update({'font.size': 14})
plt.figure(figsize=(8,6))
plt.boxplot(scores[sorted_indcs].T,labels=ftr_names[sorted_indcs],vert=False)
plt.axvline(test_score,label='test score')
plt.title("Permutation Importances (test set)")
plt.xlabel('score with perturbed feature')
plt.legend()
plt.tight_layout()
plt.show()

```



## 2.2 This is also implemented in sklearn

```
[11]: from sklearn.inspection import permutation_importance
import matplotlib.pyplot as plt

result = permutation_importance(grid, X_test, y_test,
    ↳n_repeats=10, random_state=0, scoring='neg_root_mean_squared_error')

scores = result.importances
ftr_names = X_test.columns
```

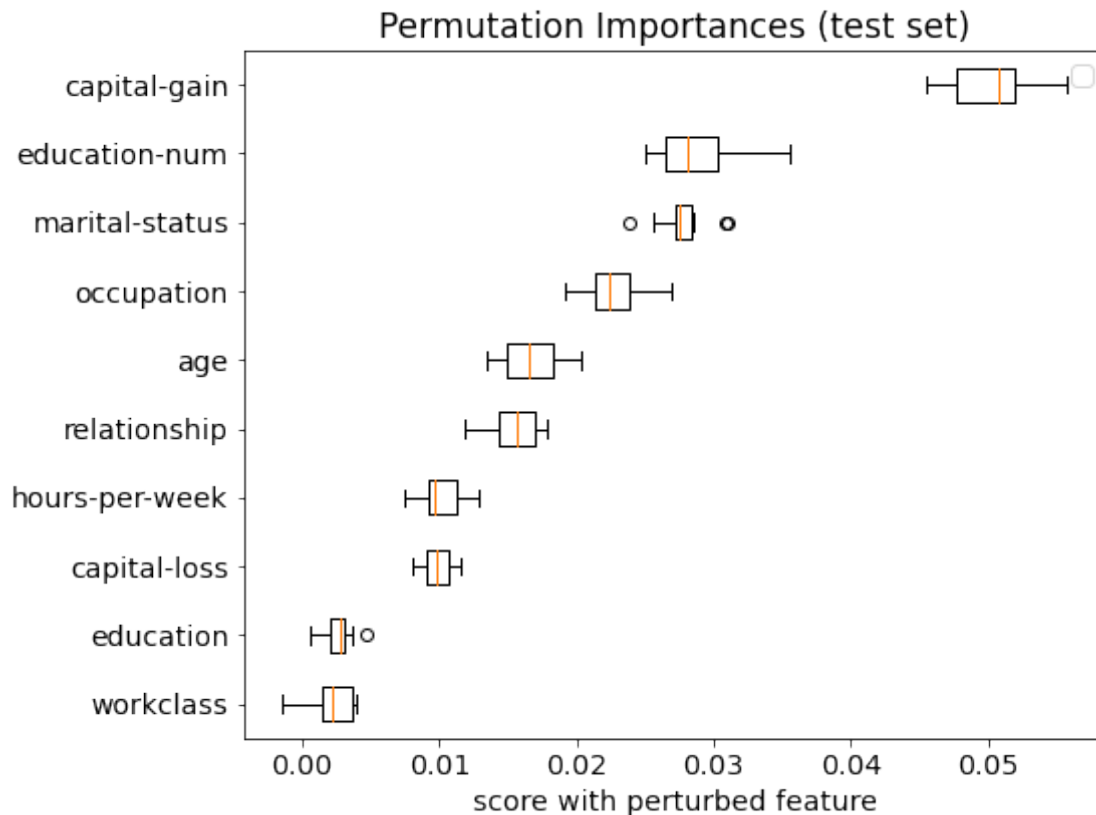
```
[12]: from sklearn.metrics import mean_squared_error
y_test_pred = grid.predict(X_test)
test_score = np.sqrt(mean_squared_error(y_test, y_test_pred))
```

```
[13]: nr_ftrs = 10

sorted_indcs = np.argsort(np.mean(scores, axis=1))
plt.rcParams.update({'font.size': 14})
plt.figure(figsize=(8,6))
```

```
plt.boxplot(scores[sorted_indcs[-nr_ftrs:]].
↳T,labels=ftr_names[sorted_indcs[-nr_ftrs:]],vert=False)
plt.title("Permutation Importances (test set)")
plt.xlabel('score with perturbed feature')
plt.legend()
plt.tight_layout()
plt.show()
```

No handles with labels found to put in legend.



## 2.3 Cons of permutation feature importance

- strongly correlated features
  - if one of the features is shuffled, the model can still use the other correlated feature
  - both features appear to be less important but they might actually be important
  - solution:
    - \* check the correlation matrix plot
    - \* remove all but one of the strongly correlated features
- no feature interactions
  - one feature might appear unimportant but combined with another feature could be important

- solution:
  - \* permute two features to measure how important feature pairs are
  - \* this can be computationally expensive

## 2.4 Global feature importance in non-linear models

- SVM:
  - SVC.coef\_\_ and SVR.coef\_\_ can be used as a metric of feature importance if **all** features are standardized
  - for linear SVMs only!
- random forest:
  - RandomForestRegressor.feature\_importances\_\_ and RandomForestClassifier.feature\_importances\_\_
  - gini importance or mean decrease impurity, see [here](#) and [here](#)
- XGBoost:
  - five different metrics are implemented, see [here](#) and [here](#)

## 3 Module 3: Local feature importance metrics

### 3.0.1 Learning objectives of this module:

- Describe the intuition behind SHAP values
- Create force, dependence, and summary plots to aid local interpretability

### 3.1 SHAP values

- one way to calculate local feature importances
- it is based on Shapely values from game theory
- read more [here](#), [here](#), and [here](#)

#### 3.1.1 Cooperative game theory

- A set of  $m$  players in a coalition generate a surplus.
- Some players contribute more to the coalition than others (different bargaining powers).
- How important is each player to the coalition?
- How should the surplus be divided fairly amongst the players?

#### 3.1.2 Cooperative game theory applied to feature attribution

- A set of  $m$  features in a model generate a prediction.
- Some features contribute more to the model than others (different predictive powers).
- How important is each feature to the model?
- How should the prediction be divided amongst the features?

#### 3.1.3 How is it calculated?

$$3.1.4 \quad \Phi_i = \sum_{S \subseteq M \setminus i} \frac{|S|!(M-|S|-1)!}{M!} [f_x(S \cup i) - f_x(S)]$$

- $\Phi_i$  - the contribution of feature  $i$
- $M$  - the number of features

- $S$  - a set of features excluding  $i$ , a vector of 0s and 1s (0 if a feature is missing)
- $|S|$  - the number of features in  $S$
- $f_x(S)$  - the prediction of the model with features  $S$

### 3.1.5 How is it calculated?

$$3.1.6 \quad \Phi_i = \sum_{S \subseteq M \setminus i} \frac{|S|!(M-|S|-1)!}{M!} [f_x(S \cup i) - f_x(S)]$$

- the difference feature  $i$  makes in the prediction:
  - $f_x(S \cup i)$  - the prediction with feature  $i$
  - $f_x(S)$  - the prediction without feature  $i$
- loop through all possible ways a set of  $S$  features can be selected from the  $M$  features excluding  $i$
- weight the contribution based on how many ways we can select  $|S|$  features

```
[14]: import numpy as np
import pandas as pd
import xgboost
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
import matplotlib.pyplot as plt

df = pd.read_csv('data/adult_data.csv')
label = 'gross-income'
y = LabelEncoder().fit_transform(df[label])
df.drop(columns=[label], inplace=True)
X = df
ftr_names = X.columns
print(X.head())
print(y)
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	

2	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female

	capital-gain	capital-loss	hours-per-week	native-country
0	2174	0	40	United-States
1	0	0	13	United-States
2	0	0	40	United-States
3	0	0	40	United-States
4	0	0	40	Cuba

[0 0 0 ... 0 0 1]

```
[15]: def ML_pipeline_kfold(X,y,random_state,n_folds):
    # create a test set
    X_other, X_test, y_other, y_test = train_test_split(X, y, test_size=0.2,
    random_state = random_state)
    # splitter for _other
    kf =
    StratifiedKFold(n_splits=n_folds,shuffle=True,random_state=random_state)
    # create the pipeline: preprocessor + supervised ML method
    cat_ftrs =
    ['workclass','education','marital-status','occupation','relationship','race','sex','native-
    cont_ftrs =
    ['age','fnlwgt','education-num','capital-gain','capital-loss','hours-per-week']
    # one-hot encoder
    categorical_transformer = Pipeline(steps=[
        ('onehot', OneHotEncoder(sparse=False,handle_unknown='ignore'))])
    # standard scaler
    numeric_transformer = Pipeline(steps=[
        ('scaler', StandardScaler())])
    preprocessor = ColumnTransformer(
        transformers=[
            ('num', numeric_transformer, cont_ftrs),
            ('cat', categorical_transformer, cat_ftrs)])
    pipe = make_pipeline(preprocessor,RandomForestClassifier(n_estimators =
    100,random_state=random_state))
    # the parameter(s) we want to tune
    param_grid = {'randomforestclassifier__max_depth': [10,30,100,300],
        'randomforestclassifier__min_samples_split': [16, 32, 64,
    128]}
    # prepare gridsearch
    grid = GridSearchCV(pipe, param_grid=param_grid,cv=kf, return_train_score =
    True,n_jobs=-1,verbose=10)
    # do kfold CV on _other
    grid.fit(X_other, y_other)
    feature_names = cont_ftrs + \
```



```

        list(grid.best_estimator_[0].named_transformers_['cat'][0].
↪get_feature_names(cat_ftrs))
    return grid, np.array(feature_names), X_test, y_test

```

```

[16]: grid, feature_names, X_test, y_test = ML_pipeline_kfold(X,y,42,4)
print(grid.best_score_)
print(grid.score(X_test,y_test))
print(grid.best_params_)

```

Fitting 4 folds for each of 16 candidates, totalling 64 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done   2 tasks      | elapsed:    5.7s
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:    9.6s
[Parallel(n_jobs=-1)]: Done  16 tasks      | elapsed:   19.7s
[Parallel(n_jobs=-1)]: Done  25 tasks      | elapsed:   19.3s
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:   24.5s
[Parallel(n_jobs=-1)]: Done  45 tasks      | elapsed:   29.5s
[Parallel(n_jobs=-1)]: Done  56 out of  64 | elapsed:   34.9s remaining:    5.0s
[Parallel(n_jobs=-1)]: Done  64 out of  64 | elapsed:   39.8s finished

0.862906941031941
0.8667280822969445
{'randomforestclassifier__max_depth': 100,
 'randomforestclassifier__min_samples_split': 64}

```

```

[17]: import shap
shap.initjs() # required for visualizations later on
# create the explainer object with the random forest model
explainer = shap.TreeExplainer(grid.best_estimator_[1])
# transform the test set
X_test_transformed = grid.best_estimator_[0].transform(X_test)
print(np.shape(X_test_transformed))
# calculate shap values on the first 1000 points in the test
shap_values = explainer.shap_values(X_test_transformed[:
↪1000],check_additivity=False)
print(np.shape(shap_values))

```

<IPython.core.display.HTML object>

Setting feature\_perturbation = "tree\_path\_dependent" because no background data was given.

```

(6513, 108)
(2, 1000, 108)

```

### 3.1.7 Explain a point


```
[18]: index = 1 # the index of the point to explain
      print(explainer.expected_value[0]) # we explain class 0 predictions
      shap.force_plot(explainer.expected_value[0], shap_values[0][index,:], features_
      ↪= X_test_transformed[index,:], feature_names = feature_names)
```

0.7589753531941029

[18]: <IPython.core.display.HTML object>


### 3.1.8 Feature value vs. shap value

```
[19]: import matplotlib
      matplotlib.rcParams.update({'font.size': 20})
      ftr = 'education-num'
      indx = np.argwhere(feature_names=='education-num')
      plt.scatter(X_test_transformed[:1000,indx],shap_values[1][:,indx])
      plt.ylabel('shap value')
      plt.xlabel(ftr)
      plt.show()
```



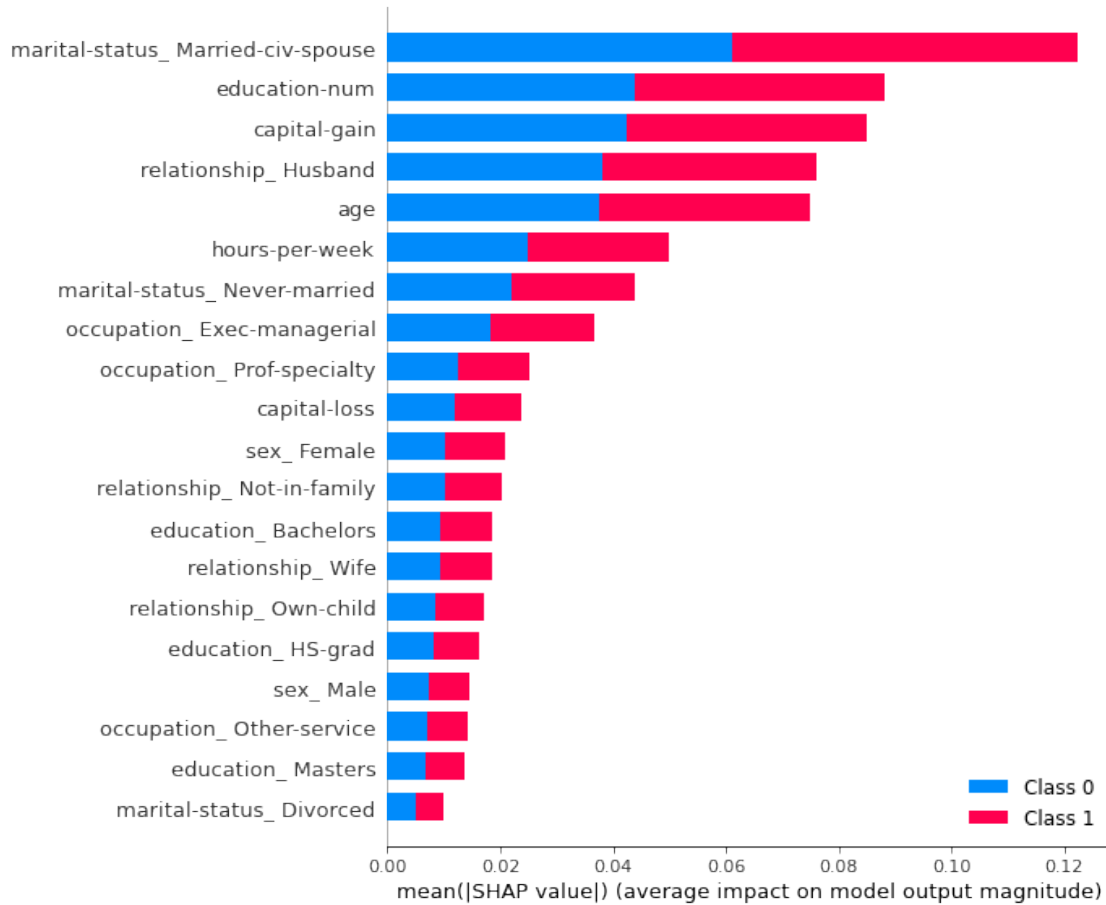
### 3.1.9 Dependence plot

```
[20]: shap.dependence_plot(ftr, shap_values[1], X_test_transformed[:1000],  
    ↪ feature_names=feature_names)
```



### 3.1.10 It can also be used for global feature importance

```
[21]: shap.summary_plot(shap_values, X_test_transformed[:1000], feature_names =  
    ↪ feature_names)
```



### 3.2 SHAP cons

- it can be numerically expensive
  - an efficient shap method was developed for trees, see [here](#)
- how to estimate  $f_x(S)$ ?
  - this is not trivial because models cannot change the number of features they use
  - usually the values of the dropped features are replaced with the mean or 0
  - this is approximate but no one came up with a better way

[ ]: